

1. Collate the 2 excel files to have all the information at one place.  
Check for missing values and duplicates before joining the 2 datasets.

```
!pip install xgboost
```

```
Defaulting to user installation because normal site-packages is not writeable
```

```
Requirement already satisfied: xgboost in c:\users\naven\appdata\roaming\python\python312\site-packages (2.1.4)
```

```
Requirement already satisfied: numpy in c:\programdata\anaconda3\lib\site-packages (from xgboost) (1.26.4)
```

```
Requirement already satisfied: scipy in c:\programdata\anaconda3\lib\site-packages (from xgboost) (1.13.1)
```

```
import xgboost as xgb  
print(xgb.__version__)
```

```
2.1.4
```

```
# All libraries/functions required
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.preprocessing import OrdinalEncoder  
import plotly.express as px  
from statsmodels.formula.api import ols  
import statsmodels.api as sm  
import scipy.stats as stats  
from statsmodels.stats.multicomp import pairwise_tukeyhsd  
from sklearn.linear_model import SGDRegressor, Ridge  
from sklearn.model_selection import KFold, StratifiedKFold,  
RandomizedSearchCV, train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import mean_squared_error as mse, r2_score  
from sklearn.preprocessing import StandardScaler  
from sklearn.ensemble import RandomForestRegressor  
import xgboost as xgb  
from sklearn.pipeline import Pipeline  
from xgboost import XGBRegressor  
from sklearn.model_selection import GridSearchCV
```

```
address = ''
```

```
hosp = pd.read_csv(r'I:\December 2024\PGC DS Capstone\Project\  
Datasets_updated\Capstone_1\Hospitalisation details.csv')  
medic = pd.read_csv(r'I:\December 2024\PGC DS Capstone\Project\  
Datasets_updated\Capstone_1\Medical Examinations.csv')
```

```
names = pd.read_excel(r'I:\December 2024\PGC DS Capstone\Project\
Datasets_updated\Capstone_1\Names.xlsx')
```

```
# hosp = pd.read_excel('/content/drive/MyDrive/Colab
Notebooks/Hospitalisation details.xlsx')
# medic = pd.read_excel('/content/drive/MyDrive/Colab
Notebooks/Medical Examinations.xlsx')
```

## Data inspection using .info()

```
hosp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2343 entries, 0 to 2342
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Customer ID     2343 non-null   object
1   year            2343 non-null   object
2   month           2343 non-null   object
3   date            2343 non-null   int64
4   children        2343 non-null   int64
5   charges         2343 non-null   float64
6   Hospital tier   2343 non-null   object
7   City tier        2343 non-null   object
8   State ID        2343 non-null   object
dtypes: float64(1), int64(2), object(6)
memory usage: 164.9+ KB
```

```
medic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 8 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   Customer ID                 2335 non-null   object
1   BMI                         2335 non-null   float64
2   HBA1C                       2335 non-null   float64
3   Heart Issues                2335 non-null   object
4   Any Transplants             2335 non-null   object
5   Cancer history              2335 non-null   object
6   NumberOfMajorSurgeries      2335 non-null   object
7   smoker                      2335 non-null   object
dtypes: float64(2), object(6)
memory usage: 146.1+ KB
```

```
names.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Customer ID  2335 non-null   object
1   name         2335 non-null   object
dtypes: object(2)
memory usage: 36.6+ KB
```

```
master_data = pd.merge(hosp, medic, how = 'inner', on = 'Customer ID')
```

```
master_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Customer ID  2335 non-null   object
1   year         2335 non-null   object
2   month        2335 non-null   object
3   date         2335 non-null   int64
4   children     2335 non-null   int64
5   charges      2335 non-null   float64
6   Hospital tier  2335 non-null   object
7   City tier     2335 non-null   object
8   State ID     2335 non-null   object
9   BMI          2335 non-null   float64
10  HBA1C        2335 non-null   float64
11  Heart Issues  2335 non-null   object
12  Any Transplants  2335 non-null   object
13  Cancer history  2335 non-null   object
14  NumberOfMajorSurgeries  2335 non-null   object
15  smoker       2335 non-null   object
dtypes: float64(3), int64(2), object(11)
memory usage: 292.0+ KB
```

```
master_data = master_data.merge(names,on='Customer ID')
```

```
master_data.head(2)
```

	Customer ID	year	month	date	children	charges	Hospital tier	City tier
0	Id2335	1992	Jul	9	0	563.84	tier - 2	tier - 3
1	Id2334	1992	Nov	30	0	570.62	tier - 2	tier - 1

	State ID	BMI	HBA1C	Heart Issues	Any Transplants	Cancer history
--	----------	-----	-------	--------------	-----------------	----------------

0	R1013	17.58	4.51	No	No	No
1	R1013	17.60	4.39	No	No	No

	NumberOfMajorSurgeries	smoker	name
0	1	No	German, Mr. Aaron K
1	1	No	Rosendahl, Mr. Evan P

```
master_data["Customer ID"].head()
```

```
0    Id2335
1    Id2334
2    Id2333
3    Id2332
4    Id2331
```

```
Name: Customer ID, dtype: object
```

```
master_data.columns = master_data.columns.str.lower()
master_data.columns = master_data.columns.str.replace(' ', '_')
master_data.columns
```

```
Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',
       'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hbalc',
       'heart_issues', 'any_transplants', 'cancer_history',
       'numberofmajorsurgeries', 'smoker', 'name'],
      dtype='object')
```

```
master_data.customer_id
```

```
0    Id2335
1    Id2334
2    Id2333
3    Id2332
4    Id2331
```

```
...
2330    Id5
2331    Id4
2332    Id3
2333    Id2
2334    Id1
```

```
Name: customer_id, Length: 2335, dtype: object
```

2. The data seems to have trivial values in a few variables. These are "?". Find the percentage of rows which have such value ("?",) in any column. Delete such rows in case you don't lose significant information.

```
(master_data == '?').sum()
```

customer_id	0
year	2
month	3
date	0
children	0
charges	0
hospital_tier	1
city_tier	1
state_id	2
bmi	0
hbalc	0
heart_issues	0
any_transplants	0
cancer_history	0
numberofmajorsurgeries	0
smoker	2
name	0

dtype: int64

*# replacing '?' with np.NA for easy access and removal*

```
miss_perc = (master_data == '?').sum(axis = 1)/master_data.shape[1] *
100
miss_perc[miss_perc > 0]
```

11	5.882353
13	5.882353
17	11.764706
542	5.882353
1046	5.882353
1049	5.882353
1700	5.882353
1775	5.882353
2165	5.882353
2332	5.882353

dtype: float64

```
miss_perc[miss_perc>0].index
```

```
Index([11, 13, 17, 542, 1046, 1049, 1700, 1775, 2165, 2332],
      dtype='int64')
```

```
miss_perc_col = (master_data == '?').sum(axis =
0)/master_data.shape[0] * 100
miss_perc_col.sort_values(ascending= False)
```

month	0.128480
state_id	0.085653
smoker	0.085653
year	0.085653
hospital_tier	0.042827

```

city_tier          0.042827
heart_issues       0.000000
numberofmajorsurgeries 0.000000
cancer_history     0.000000
any_transplants    0.000000
customer_id        0.000000
hbalc              0.000000
bmi                0.000000
charges            0.000000
children           0.000000
date               0.000000
name               0.000000
dtype: float64

```

```

master_noq = master_data.drop(index = miss_perc[miss_perc>0].index)
master_noq.shape

```

```

(2325, 17)

```

```

master_noq.isna().sum()

```

```

customer_id        0
year                0
month              0
date               0
children           0
charges            0
hospital_tier       0
city_tier           0
state_id           0
bmi                0
hbalc              0
heart_issues        0
any_transplants     0
cancer_history      0
numberofmajorsurgeries 0
smoker             0
name               0
dtype: int64

```

3. The data has nominal and ordinal categorical variables. How are you going to incorporate these variables in the next steps of modelling. Use necessary transformation methods to deal with nominal and ordinal types of data.

```

master_noq[['city_tier', 'hospital_tier']]

```

```

   city_tier hospital_tier
0   tier - 3   tier - 2

```

```

1      tier - 1      tier - 2
2      tier - 1      tier - 2
3      tier - 3      tier - 3
4      tier - 3      tier - 3
...
2329   tier - 3      tier - 1
2330   tier - 2      tier - 1
2331   tier - 3      tier - 1
2333   tier - 3      tier - 2
2334   tier - 3      tier - 1

```

```
[2325 rows x 2 columns]
```

```
master_noq.state_id.value_counts()
```

```
state_id
```

```

R1013    609
R1011    574
R1012    572
R1024    159
R1026     84
R1021     70
R1016     64
R1025     40
R1023     38
R1017     36
R1019     26
R1022     14
R1014     13
R1015     11
R1018      9
R1020      6

```

```
Name: count, dtype: int64
```

```
master_noq.state_id.head()
```

```

0      R1013
1      R1013
2      R1013
3      R1013
4      R1013

```

```
Name: state_id, dtype: object
```

```
master_noq.state_id.nunique()
```

```
16
```

```
# Using ordinalencoder to deal with ordinal categorical variables -  
city tier and hospital tier
```

```
ordinal = OrdinalEncoder(categories= [['tier - 3', 'tier - 2', 'tier -
```

```
1'], ['tier - 3', 'tier - 2', 'tier - 1'])
master_noq[['city_tier_ord', 'hospital_tier_ord']] =
ordinal.fit_transform(master_noq[['city_tier', 'hospital_tier']])
```

```
pd.crosstab(master_noq['city_tier_ord'], master_noq['city_tier'])
```

city_tier	tier - 1	tier - 2	tier - 3
city_tier_ord			
0.0	0	0	789
1.0	0	807	0
2.0	729	0	0

```
pd.crosstab(master_noq['hospital_tier_ord'], master_noq['hospital_tier'])
```

hospital_tier	tier - 1	tier - 2	tier - 3
hospital_tier_ord			
0.0	0	0	691
1.0	0	1334	0
2.0	300	0	0

```
master_noq.head(3)
```

	customer_id	year	month	date	children	charges	hospital_tier
city_tier \							
0	Id2335	1992	Jul	9	0	563.84	tier - 2 tier - 3
1	Id2334	1992	Nov	30	0	570.62	tier - 2 tier - 1
2	Id2333	1993	Jun	30	0	600.00	tier - 2 tier - 1

	state_id	bmi	hbalc	heart_issues	any_transplants
cancer_history \					
0	R1013	17.58	4.51	No	No No
1	R1013	17.60	4.39	No	No No
2	R1013	16.47	6.35	No	No Yes

	numberofmajorsurgeries	smoker	name	city_tier_ord
\				
0	1	No	German, Mr. Aaron K	0.0
1	1	No	Rosendahl, Mr. Evan P	2.0
2	1	No	Albano, Ms. Julie	2.0

```
hospital_tier_ord
```



0	1.0
1	1.0
2	1.0

4. State ID has around 16 states. The data does not have proportional representation of all the states. Also creating dummy variables corresponding to all the regions may lead to too many insignificant predictors. Nevertheless, only R1011, R1012 and R1013 are important to look deeper into. Keeping these ideas in mind, come up with a suitable strategy here.

```
vc = master_noq.state_id.value_counts()    # frequency of each category
vc

state_id
R1013    609
R1011    574
R1012    572
R1024    159
R1026     84
R1021     70
R1016     64
R1025     40
R1023     38
R1017     36
R1019     26
R1022     14
R1014     13
R1015     11
R1018      9
R1020      6
Name: count, dtype: int64

vc[:3].index                                # picking top 3 most
frequent categories

Index(['R1013', 'R1011', 'R1012'], dtype='object', name='state_id')

for i in vc[:3].index:
    var_name = 'state_id_' + i    # create name for the dummy variable
    print(var_name)
    master_noq[var_name] = 0      # giving a dummy value 0 to dummy
variable
    master_noq.loc[master_noq.state_id == i, var_name] = 1 # replacing
0 by 1 where state id is equal to category of the dummy variable
```

```

state_id_R1013
state_id_R1011
state_id_R1012

master_noq.state_id.value_counts()

state_id
R1013    609
R1011    574
R1012    572
R1024    159
R1026     84
R1021     70
R1016     64
R1025     40
R1023     38
R1017     36
R1019     26
R1022     14
R1014     13
R1015     11
R1018      9
R1020      6
Name: count, dtype: int64

# checking the no of records corresponding to R1013

master_noq['state_id_R1013'].value_counts()

state_id_R1013
0    1716
1     609
Name: count, dtype: int64

master_noq['state_id_R1012'].value_counts()

state_id_R1012
0    1753
1     572
Name: count, dtype: int64

```

5. Variable 'NumberOfMajorSurvalue\_counts seems to have string values as well. You may want to clean this variable.

```

master_noq.numberofmajorsurgeries.unique()

array(['1', 'No major surgery', '2', '3'], dtype=object)

master_noq.loc[master_noq.numberofmajorsurgeries == 'No major
surgery', 'numberofmajorsurgeries' ] = 0

```

```

master_noq.numberofmajorsurgeries =
master_noq.numberofmajorsurgeries.astype(int)

master_noq.numberofmajorsurgeries.unique()

array([1, 0, 2, 3])

```

6. Age seems to be an important factor for this analysis. Based on date of birth information, calculate the age of the patients.

```

master_noq.year = master_noq.year.astype(int)
master_noq['age'] = 2025 - master_noq.year
master_noq.head(2)

```

	customer_id	year	month	date	children	charges	hospital_tier
0	Id2335	1992	Jul	9	0	563.84	tier - 2
1	Id2334	1992	Nov	30	0	570.62	tier - 2

	state_id	bmi	...	cancer_history	numberofmajorsurgeries
0	R1013	17.58	...	No	1
1	R1013	17.60	...	No	1

	state_id_R1013	name	city_tier_ord	hospital_tier_ord
0	German, Mr.	Aaron K	0.0	1.0
1	Rosendahl, Mr.	Evan P	2.0	1.0

	state_id_R1011	state_id_R1012	age
0	0	0	33
1	0	0	33

[2 rows x 23 columns]

7. Gender of the patient may be an important factor to decide the hospitalization cost. Salutation provided in the name of the beneficiary can be used to determine the gender. Create a new field for the gender of beneficiary.

```

master_noq.name

```

```

0          German, Mr.  Aaron K
1      Rosendahl, Mr.  Evan P
2          Albano, Ms.  Julie
3      Riveros Gonzalez, Mr.  Juan D. Sr.
4          Brietzke, Mr.  Jordan

```

```

...
2329      Baker, Mr.  Russell B.
2330      Kadala, Ms.  Kristyn
2331      Osborne, Ms.  Kelsey
2333      Lehner, Mr.  Matthew D
2334      Hawks, Ms.  Kelly

```

```
Name: name, Length: 2325, dtype: object
```

```

master_noq['title'] =
master_noq.name.str.split('[,.]').str[1].str.strip()

```

```
master_noq.title.value_counts()
```

```

title
Mr      1160
Ms      1023
Mrs     142
Name: count, dtype: int64

```

```
master_noq.shape
```

```
(2325, 24)
```

```
1160+1023+142
```

```
2325
```

```

master_noq['gender'] = 'female'
master_noq.loc[master_noq.title == 'Mr', 'gender'] = 'male'
master_noq.loc[master_noq.title == 'Mrs']

```

	customer_id	year	month	date	children	charges	
hospital_tier \							
24	Id2311	2001	Aug	19	0	964.71	tier - 3
172	Id2163	2004	Dec	27	0	1863.45	tier - 3
197	Id2138	2004	Jun	12	0	2094.10	tier - 3
328	Id2007	1993	Sep	25	0	3162.02	tier - 2
348	Id1987	2003	Dec	5	0	3300.70	tier - 2
...	...	...	...	...	...	...	...
1790	Id545	1963	Jul	4	0	18208.34	tier - 1

1808	Id527	1963	Dec	6	0	18883.33	tier - 1
1811	Id524	1963	Oct	20	0	18954.56	tier - 1
1839	Id496	1966	Aug	10	0	19995.29	tier - 1
1848	Id487	1962	Jul	2	0	20354.50	tier - 3

	city_tier	state_id	bmi	...	smoker	
name \						
24	tier - 2	R1013	25.19	...	No	Keys, Mrs.
Kathleen						
172	tier - 1	R1025	27.06	...	No	Stanislav, Mrs. Grace
H						
197	tier - 2	R1025	27.74	...	No	Padula, Mrs.
Lauren						
328	tier - 3	R1013	25.61	...	No	Martin, Mrs. Kristen
M						
348	tier - 2	R1025	30.54	...	No	Mendez-Karr, Mrs.
Cynthia						
...	...	...	...	...	...	..
.						
1790	tier - 2	R1026	44.20	...	No	Shigezumi, Mrs.
Teiko						
1808	tier - 1	R1026	46.19	...	No	Hughey, Mrs. Ashley
E						
1811	tier - 1	R1026	46.40	...	No	Rogers, Mrs. Anita
L.						
1839	tier - 3	R1026	51.74	...	No	Oehlke, Mrs.
Jessica						
1848	tier - 2	R1026	49.77	...	No	Argall, Mrs. Tara
R						

	city_tier_ord	hospital_tier_ord	state_id_R1013	state_id_R1011	\
24	1.0	0.0	1	0	
172	2.0	0.0	0	0	
197	1.0	0.0	0	0	
328	0.0	1.0	1	0	
348	1.0	1.0	0	0	
...	...	...	...	...	
1790	1.0	2.0	0	0	
1808	2.0	2.0	0	0	
1811	2.0	2.0	0	0	
1839	0.0	2.0	0	0	
1848	1.0	0.0	0	0	

	state_id_R1012	age	title	gender
24	0	24	Mrs	female

172	0	21	Mrs	female
197	0	21	Mrs	female
328	0	32	Mrs	female
348	0	22	Mrs	female
...	...	...	...	...
1790	0	62	Mrs	female
1808	0	62	Mrs	female
1811	0	62	Mrs	female
1839	0	59	Mrs	female
1848	0	63	Mrs	female

[142 rows x 25 columns]




master\_noq['gender']

0	male
1	male
2	female
3	male
4	male
...	...
2329	male
2330	female
2331	female
2333	male
2334	female

Name: gender, Length: 2325, dtype: object


8. Visualize the distribution of cost using histogram, box and whisker and swarm plot. How the distribution is different across gender and different tiers of hospitals. Share your observation.

```
plt.figure(figsize = (25,10))
grid = plt.GridSpec(2, 2, wspace=0.4, hspace=0.3)
plt.subplot(grid[0, 0])
plt.hist(master_noq.charges, bins = 100)
plt.subplot(grid[0, 1])
master_noq.charges.plot.kde()
plt.subplot(grid[1, :])
plt.boxplot(master_noq.charges, vert = False)
plt.show()
```




## WRT gender

```
plt.figure(figsize = (15,5))
sns.boxplot(x = "charges",y = "gender", data = master_noq)
plt.show()
```




## WRT city tier

```
plt.figure(figsize = (15,5))
sns.boxplot(x = "charges",y = "city_tier", data = master_noq)
plt.show()
```



## WRT Hospital tier

```
plt.figure(figsize = (15,5))
sns.boxplot(x = "charges",y = "hospital_tier", data = master_noq)
plt.show()
```




9. Create a radar chart to showcase the median hospitalization cost across different tiers of hospitals.

```
median = master_noq.groupby('hospital_tier')
[['charges']].median().reset_index()

fig = px.line_polar(median, r='charges', theta='hospital_tier') #,
line_close=True
fig.show()
```






10. Create a frequency table and hence a stacked bar-chart to visualize the count of people in different tiers of cities and hospitals.

```
pd.crosstab(master_noq.city_tier, master_noq.hospital_tier)
```

hospital_tier	tier - 1	tier - 2	tier - 3
city_tier			
tier - 1	85	403	241
tier - 2	106	479	222
tier - 3	109	452	228

```
pd.crosstab(master_noq.city_tier,  
master_noq.hospital_tier).plot.bar(stacked = True)  
plt.show()
```



## 11. Test the following null hypotheses:

- Average hospitalization cost across the 3 types of hospitals is not significantly different
- Average hospitalization cost across the 3 types of cities is not significantly different
- Average hospitalization cost for smokers is not significantly different than non-smokers
- Smoking and Hearth issues are independent

H0 : Average hospitalization cost across the 3 types of hospitals is not significantly different

```
mod = ols('charges ~ hospital_tier', data = master_noq).fit()
res = sm.stats.anova_lm(mod)
res
```

	df	sum_sq	mean_sq	F
PR(>F)				
hospital_tier	2.0	9.763011e+10	4.881505e+10	493.989566
1.773822e-179				
Residual	2322.0	2.294554e+11	9.881799e+07	NaN
NaN				

H0 = Average hospitalization cost across the 3 types of cities is not significantly different

```
mod = ols('charges ~ city_tier', data = master_noq).fit()
res = sm.stats.anova_lm(mod)
res
```

	df	sum_sq	mean_sq	F	PR(>F)
city_tier	2.0	4.092192e+08	2.046096e+08	1.454356	0.233763
Residual	2322.0	3.266763e+11	1.406874e+08	NaN	NaN

H0: Average hospitalization cost for smokers is not significantly different than non-smokers

```
sample1 = master_noq.loc[master_noq.smoker == 'yes', 'charges']
sample2 = master_noq.loc[master_noq.smoker != 'yes', 'charges']
stats.ttest_ind(sample1, sample2)
```

```
TtestResult(statistic=74.15560699695726, pvalue=0.0, df=2323.0)
```

H0 : Smoking and Heart issues are independent

```
observed_table = pd.crosstab(master_noq.smoker,
master_data.heart_issues)
```

```
observed_table
```

heart_issues	No	yes
smoker		
No	1108	731
yes	297	189

```
chi, p, df, expected = stats.chi2_contingency(observed_table)
```

```
chi, p, df, expected
```

```
(0.08588150449910657,
0.7694797581780767,
1,
array([[1111.30967742, 727.69032258],
       [ 293.69032258, 192.30967742]]))
```

12. Check the correlation between predictors to identify highly correlated predictors. Visualize using a heatmap.

```
master_noq.columns
```

```
Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',
'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hbalc',
'heart_issues', 'any_transplants', 'cancer_history',
'numberofmajorsurgeries', 'smoker', 'name', 'city_tier_ord',
```

```

'hospital_tier_ord', 'state_id_R1013', 'state_id_R1011',
'state_id_R1012', 'age', 'title', 'gender'],
dtype='object')

```

```

data = master_noq.drop(columns = ['customer_id', 'name', 'year',
'month', 'date', 'hospital_tier',
'city_tier', 'state_id', 'title'])

```

```


corr_plot = data.select_dtypes(exclude='object').corr()
ma = np.ones_like(corr_plot)
ma[np.tril_indices_from(ma)] = 0

```

```

plt.figure(figsize = (18,5))
sns.heatmap(corr_plot, annot= True , mask = ma, cmap='PuRd')
plt.show()

```



### 13. Final Model Development and Evaluation :

Perform stratified 5-fold cross validation technique for final prediction and validation. Make sure to use standardization, hyperparameter tuning effectively. There must be effective use of sklearn-pipelines.

a. Create 5 fold in the data. You may want to create a variable to identify the folds.

```

data_2 = pd.get_dummies(data, drop_first=True)
data_2.reset_index(drop=True, inplace = True)
data_2.head()

```

	children	charges	bmi	hba1c	numberofmajorsurgeries
city_tier_ord \					
0	0	563.84	17.58	4.51	1
0.0					
1	0	570.62	17.60	4.39	1
2.0					

2	0	600.00	16.47	6.35	1
2.0					
3	0	604.54	17.70	6.28	1
0.0					
4	0	637.26	22.34	5.57	1
0.0					

	hospital_tier_ord	state_id_R1013	state_id_R1011	state_id_R1012
age \				
0	1.0	1	0	0
33				
1	1.0	1	0	0
33				
2	1.0	1	0	0
32				
3	0.0	1	0	0
33				
4	0.0	1	0	0
27				

	heart_issues_yes	any_transplants_yes	cancer_history_Yes
smoker_yes \			
0	False	False	False
False			
1	False	False	False
False			
2	False	False	True
False			
3	False	False	False
False			
4	False	False	False
False			

	gender_male
0	True
1	True
2	False
3	True
4	True

```
# rearrange data to put 'charges' as first column or last
model_data = data_2.drop(columns = 'charges')
model_data.head()
model_data['charges'] = data_2.charges
model_data.head()
```

	children	bmi	hba1c	numberofmajorsurgeries	city_tier_ord \
0	0	17.58	4.51	1	0.0
1	0	17.60	4.39	1	2.0
2	0	16.47	6.35	1	2.0

3	0	17.70	6.28	1	0.0
4	0	22.34	5.57	1	0.0

	hospital_tier_ord	state_id_R1013	state_id_R1011	state_id_R1012
age \				
0	1.0	1	0	0
33				
1	1.0	1	0	0
33				
2	1.0	1	0	0
32				
3	0.0	1	0	0
33				
4	0.0	1	0	0
27				

	heart_issues_yes	any_transplants_yes	cancer_history_Yes
smoker_yes \			
0	False	False	False
False			
1	False	False	False
False			
2	False	False	True
False			
3	False	False	False
False			
4	False	False	False
False			

	gender_male	charges
0	True	563.84
1	True	570.62
2	False	600.00
3	True	604.54
4	True	637.26

```
model_data.columns = model_data.columns.str.lower()
```

```
model_data.columns
```

```
Index(['children', 'bmi', 'hbalc', 'numberofmajorsurgeries',
      'city_tier_ord',
      'hospital_tier_ord', 'state_id_r1013', 'state_id_r1011',
      'state_id_r1012', 'age', 'heart_issues_yes',
      'any_transplants_yes',
      'cancer_history_yes', 'smoker_yes', 'gender_male', 'charges'],
      dtype='object')
```

```
# converting y to categorical for stratified k fold
```

```
y = model_data['charges']
```

```
X = model_data.drop(columns = 'charges')
```

```
X.head()
```

	children	bmi	hbalc	numberofmajorsurgeries	city_tier_ord \
0	0	17.58	4.51	1	0.0
1	0	17.60	4.39	1	2.0
2	0	16.47	6.35	1	2.0
3	0	17.70	6.28	1	0.0
4	0	22.34	5.57	1	0.0

	hospital_tier_ord	state_id_r1013	state_id_r1011	state_id_r1012
age \				
0	1.0	1	0	0
33				
1	1.0	1	0	0
33				
2	1.0	1	0	0
32				
3	0.0	1	0	0
33				
4	0.0	1	0	0
27				

	heart_issues_yes	any_transplants_yes	cancer_history_yes
smoker_yes \			
0	False	False	False
False			
1	False	False	False
False			
2	False	False	True
False			
3	False	False	False
False			
4	False	False	False
False			

	gender_male
0	True
1	True
2	False
3	True
4	True

```
#scoring='neg_root_mean_squared_error'
```

```

#Setting up a pipeline
pipeline = Pipeline(steps=[('scaler', StandardScaler()), ('regressor',
Ridge())])

# Defining the parameters for hyperparameter tuning
parameters = {'regressor__alpha': [0.001, 0.01, 0.1, 1, 10, 100]}

# Creating the KFold object
kfold = KFold(n_splits=5, shuffle=True, random_state=42)

# Creating the grid search object
model_ridge = GridSearchCV(pipeline, parameters, cv=kfold,
scoring='neg_mean_squared_error')

model_ridge.fit(X, y)

GridSearchCV(cv=KFold(n_splits=5, random_state=42, shuffle=True),
            estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                      ('regressor', Ridge())]),
            param_grid={'regressor__alpha': [0.001, 0.01, 0.1, 1, 10,
100]},
            scoring='neg_mean_squared_error')

# Getting the best parameters and the best model
model_ridge.best_params_
{'regressor__alpha': 10}

model_ridge.best_estimator_

Pipeline(steps=[('scaler', StandardScaler()), ('regressor',
Ridge(alpha=10))])

```

## Gradient Boosting Algorithm

```

from sklearn.ensemble import GradientBoostingRegressor

# Assuming df is your DataFrame
# Use df appropriately to prepare X (input) and y (output)

# Split the data into training and testing sets
# (Make sure to replace X and y with your data appropriately)
X_train,X_test,y_train,y_test = train_test_split(X,y)
# Train the XGBoost model
model = GradientBoostingRegressor()
model.fit(X_train, y_train)

# You can print the feature importances if needed
print(model.feature_importances_)

# Identify redundant variables based on the importance scores

```



```
[3.43110020e-03 1.09566693e-01 4.62959352e-03 2.44790703e-05
 0.00000000e+00 2.59776218e-02 5.73988790e-03 6.66669833e-03
 2.42889525e-04 9.32630891e-02 1.53224388e-04 1.73460830e-05
 3.26352026e-04 7.49727697e-01 2.33328128e-04]
```

## Variable Importance

```
pd.DataFrame({'Features':model.feature_names_in_, 'Importance':model.feature_importances_}).sort_values("Importance",ascending=False)
```

	Features	Importance
13	smoker_yes	0.749728
1	bmi	0.109567
9	age	0.093263
5	hospital_tier_ord	0.025978
7	state_id_r1011	0.006667
6	state_id_r1013	0.005740
2	hbalc	0.004630
0	children	0.003431
12	cancer_history_yes	0.000326
8	state_id_r1012	0.000243
14	gender_male	0.000233
10	heart_issues_yes	0.000153
3	numberofmajorsurgeries	0.000024
11	any_transplants_yes	0.000017
4	city_tier_ord	0.000000

```
# train score
```

```
model.score(X_train,y_train)
```

```
0.9380089945557816
```

```
# test score
```

```
model.score(X_test,y_test)
```

```
0.9017407518696134
```

15. Predict the hospitalization cost for Christopher, Ms. Jayna (Date of birth – 12/28/1988, height 170 cm and weight 85 kgs). She resides in a tier1 city (state : stateid = R1011) with husband and 2 of her kids. She is tested non-diabetic ( hbA1c = 5.8). She smokes but otherwise she is healthy, no transplants and no major surgeries so far. Her father had lung cancer and that was the reason of his early demise. Hospitalization cost to predicted considering tier1 hospitals.

Find predicted hospitalization cost based on all the 5 models. The predicted value should be mean of all the 5 predicted values from the 5 models.

```
model_data.columns
Index(['children', 'bmi', 'hba1c', 'numberofmajorsurgeries',
      'city_tier_ord',
      'hospital_tier_ord', 'state_id_r1013', 'state_id_r1011',
      'state_id_r1012', 'age', 'heart_issues_yes',
      'any_transplants_yes',
      'cancer_history_yes', 'smoker_yes', 'gender_male', 'charges'],
      dtype='object')
```

```
pred_data = pd.DataFrame({'Name' : ['Christopher, Ms. Jayna'],
                          'DOB' : ['12/28/1988'],
                          'city_tier' : ['tier - 1'], 'children' :[ 2],
                          'HbA1c' : [5.8],
                          'smoker_yes' : [1],
                          'heart_issues_yes' : [0],
                          'any_transplants_yes' : [0],
                          'numberofmajorsurgeries' :[ 0],
                          'cancer_history_yes' : [1],
                          'hospital_tier' : ['tier - 1'],
                          'bmi' : [85/(1.70 **2)],
                          'state_id_R1011' : [1]
                          })
```

pred\_data

	Name	DOB	city_tier	children	HbA1c
0	Christopher, Ms. Jayna	12/28/1988	tier - 1	2	5.8
1					
	heart_issues_yes	any_transplants_yes	numberofmajorsurgeries		

```

0          0          0          0
cancer_history_yes hospital_tier      bmi  state_id_R1011
0          1      tier - 1  29.411765          1

pred_data.columns = pred_data.columns.str.lower()

# we will create columns according to the final model data already created

pred_data['gender_male'] = 0
pred_data.loc[pred_data.name.str.split('[,.]').str[1] == 'Mr',
'gender_male'] = 1
pred_data.drop(columns = 'name', inplace = True)

pred_data
      dob city_tier  children  hbalc  smoker_yes  heart_issues_yes \
0  12/28/1988  tier - 1         2    5.8          1          0

      any_transplants_yes  numberofmajorsurgeries  cancer_history_yes \
0          0          0          0          1

      hospital_tier      bmi  state_id_r1011  gender_male
0      tier - 1  29.411765          1          0

#pred_data['age'] = 2023 - pred_data.dob.astype(np.datetime64).dt.year
pred_data.drop(columns = 'dob', inplace = True)

pred_data[['city_tier_ord', 'hospital_tier_ord']] =
ordinal.transform(pred_data[['city_tier', 'hospital_tier']])

pred_data.drop(columns = ['city_tier', 'hospital_tier'], inplace = True
)

# initializing the missing columns with 0 and not include charges
for col in model_data.columns:
    if col not in pred_data.columns and col != 'charges':
        pred_data[col] = 0

pred_data
      children  hbalc  smoker_yes  heart_issues_yes  any_transplants_yes \
0          2    5.8          1          0          0

      numberofmajorsurgeries  cancer_history_yes      bmi

```

```

state_id_r1011  \
0              0              1  29.411765
1

gender_male  city_tier_ord  hospital_tier_ord  state_id_r1013  \
0           0           2.0           2.0           0

state_id_r1012  age
0           0           0

```

### Apply Gradient BOOST model for predi

```
model_data.columns
```

```

Index(['children', 'bmi', 'hbalc', 'numberofmajorsurgeries',
      'city_tier_ord',
      'hospital_tier_ord', 'state_id_r1013', 'state_id_r1011',
      'state_id_r1012', 'age', 'heart_issues_yes',
      'any_transplants_yes',
      'cancer_history_yes', 'smoker_yes', 'gender_male', 'charges'],
      dtype='object')

```

```
pred_data.columns
```

```

Index(['children', 'hbalc', 'smoker_yes', 'heart_issues_yes',
      'any_transplants_yes', 'numberofmajorsurgeries',
      'cancer_history_yes',
      'bmi', 'state_id_r1011', 'gender_male', 'city_tier_ord',
      'hospital_tier_ord', 'state_id_r1013', 'state_id_r1012',
      'age'],
      dtype='object')

```

```
pred_data=pred_data[model_data.drop(columns='charges').columns]
```

```
model.predict(pred_data)
```

```
array([22677.0047916])
```

```
test =
```

```
pd.read_html("https://en.wikipedia.org/wiki/List_of_the_busiest_airports_in_the_United_States")
```

```
len(test)
```

```
10
```

```
test[4]
```

```

Rank
Rank
0    1
1    2
2    3

```

	Airport name \
	Airport name
Memphis International Airport	
Ted Stevens Anchorage International Airport	
Louisville Muhammad Ali International Airport	

3	4	O'Hare International Airport
4	5	Miami International Airport
5	6	Los Angeles International Airport
6	7	Cincinnati/Northern Kentucky International Air...
7	8	Indianapolis International Airport
8	9	Dallas/Fort Worth International Airport
9	10	Ontario International Airport

	Location	IATA code	Cargo	
	Location	IATA code	Ibs.	% chg. 2017/16
0	Memphis, Tennessee	MEM	23949525780	00.35%
1	Anchorage, Alaska	ANC	17337337377	02.79%
2	Louisville, Kentucky	SDF	13403682652	04.68%
3	Chicago, Illinois	ORD	10373559593	010.84%
4	Miami, Florida	MIA	7963988407	00.82%
5	Los Angeles, California	LAX	7197930264	03.85%
6	Hebron, Kentucky	CVG	5700282994	033.32%
7	Indianapolis, Indiana	IND	5138500318	0-3.58%
8	Irving, Texas	DFW	4155362297	07.65%
9	Ontario, California	ONT	3522510318	015.81%